How should the penalty be chosen when the UFLP QUBO is solved by simulated annealing?

A short computational study building on Li, Hennig, Kochenberger & Glover (2026), "Quantum bridge analytics: a QUBO approach to the uncapacitated facility location problem", Annals of Operations Research. All results are from classical simulated annealing; none are quantum or Gurobi results.

1. Question

The paper reformulates the uncapacitated facility location problem (UFLP) as a QUBO (eq. 15) with a penalty P. With Gurobi's QUBO solver it found that P = 25–100% of the optimal cost always gave the optimum on the ORLIB classes, and it chose 50%. I asked three questions:

1. **Does the choice of P behave the same when the QUBO is solved by single-flip simulated annealing (SA)?**
2. **Can a good P be predicted for a new instance from cheap features, without the optimum?**
3. **What changes when capacity constraints, and the slack variables they need, are added?** The paper lists general inequality constraints as future work (§5). The model is the capacitated one in Li (2026), *Urban Lifeline* 4:7, Appendix A.

2. Setup

- **Model.** Both constructions in the paper: eq. 15 built algebraically, and the Q matrix from Properties 1–2. They give the same matrix on all 37 files. Both reproduce the appendix example entry for entry, including the −710 + 1,000 = 290 result.
- **Instances.** Classes A-1, A-2 and B-1 (37 ORLIB files). Two points about this data:
  - the capacitated-set files (cap4x, 5x, 6x, 8x, 9x, 11x, 12x) have the same fixed and service costs as cap7x, 10x and 13x, so the 37 files are **12 distinct uncapacitated problems**; I average over distinct problems;
  - an exact MILP solve (HiGHS) reproduces all 37 published optima, and **the LP relaxation's optimal solution is already 0/1 on every file**, so the LP bound equals the optimum.
- **Solver and grid.** SA from `dwave-samplers` (50 reads × 1,000 sweeps, default schedule), 10 seeds per file. P = r × optimum, with r ∈ {1, 2, 5, 10, **25, 50, 75, 100**}% (the paper's values in bold); 2,960 runs in total.
- **Quality.** Gap = (cost − opt)/opt. I report it two ways:
  - **raw:** the best feasible read, as in the paper;
  - **repaired:** keep the facilities a read opens and assign each customer to its cheapest open one.

3. Findings

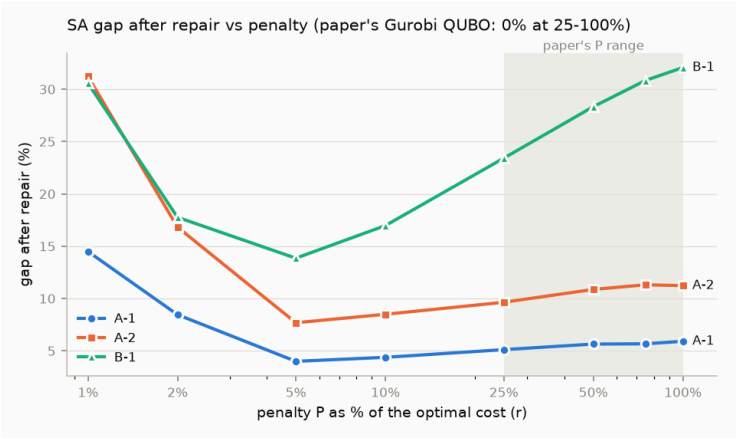
(a) Why the paper saw no infeasibility at 25%. A short argument gives bounds on the smallest "exact" P, meaning the smallest P at which every minimizer of the QUBO is feasible, and hence optimal:

- *Sufficient:* $P > \max(\max_j \min_i (f_i + c_{ij}), \max_i f_i)$. Above this, every infeasible state has a strictly improving move: serve a dropped customer, remove a duplicate assignment, or open a facility that is serving customers.
- *Necessary:* $P \geq \max_j \min_i c_{ij}$, and $P \geq \text{opt}/|J|$.

On the toy example the exact threshold is 72.5 = opt/|J|. On the 12 ORLIB problems the sufficient bound is **16.6–22.6% of the optimum**, so every QUBO in the paper's 25–100% range was provably exact. SA agrees: at 25–100% the best read of all 1,480 runs is feasible, while no read below 25% is feasible.

(b) For SA, the paper's range is the wrong place to look.

Mean gap % (12 problems)	A-1	A-2	B-1
raw, at P = 50% (as measured in the paper)	56.6	80.1	109.6
repaired, at P = 50%	5.6	10.9	28.3
repaired, at **P = 5%** (SA's best)	**4.0**	**7.7**	**13.9**
paper, Gurobi QUBO solver, P = 25–100%	0	0	0



- The repaired gap is lowest at P = 5% for 10 of the 12 problems. That is below the exactness threshold: SA's raw reads are infeasible there, but the facilities they open are good.
- Across 25–100%, SA gets steadily worse as P grows, whereas Gurobi's solution quality did not depend on P on these classes.
- **More effort does not fix this.** Ten times more sweeps on cap74 improved the repaired gap by at most 1.6 points, and not at all at P ≥ 25%. This is consistent with energy barriers of size about P: moving one customer between facilities passes through a state that costs about P.
- **What makes an instance hard differs by solver.** Although every file has a zero LP gap, SA struggles most where fixed costs are high and few facilities are open (cap74, cap104 and cap134, with 4 open). On B-1 at P = 50%, the gap ranges from 13.6% (cap131) to 45.7% (cap134).

(c) A good penalty depends on the solver, and here on the instance too. I generated 200 seeded instances in four families, 16–50 sites × 50 customers. Three resemble ORLIB (LP gap ≈ 0). The fourth is modelled on the Kochetov–Ivanenko construction behind class B-2, and has LP gaps of 1.3–26%. Each instance was solved exactly, then by SA at 12 values of P = r × LP bound (7,200 runs). The best r is always 2–6%, but fixed rules fail unevenly:

Mean repaired gap %	best r per instance	fixed 3%	fixed 4%	50% of LP	100% of LP
all 200 instances	1.30	12.22	2.94	42.3	48.7
large-gap family only	2.29	45.18	6.70	46.0	51.9

A fixed 3% is nearly optimal on the three ORLIB-like families but costs 45% on the large-gap family. I trained a gradient-boosted regressor on log r* using ten scale-free features, none of which uses the optimum: the greedy–LP gap, the exactness bounds relative to LP, cost dispersion, and others. The results are mixed:

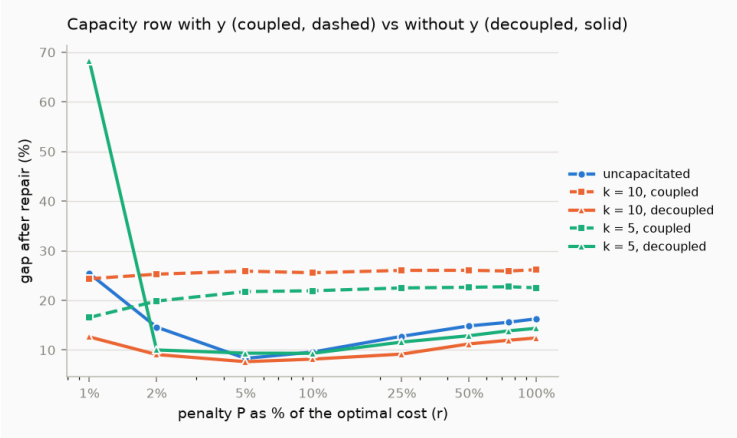
- **5-fold cross-validation:** a mean gap of 1.69%, against 2.94% for the best fixed r (paired 95% CI of the difference [0.86, 1.67]) and 1.30% for the oracle.
- **Large-gap family held out from training:** 6.7% against 45.2% for the best fixed rule. A simple rule, r = k × (exactness bound)/LP, does about as well here.
- **ORLIB:** it does not transfer; it is 1.8 points worse than a fixed 4%. Its predictions mostly separate families rather than instances.
- **Why not ORLIB.** ORLIB's demands are heavy-tailed: one customer holds 22% of all demand, which put ORLIB outside the training data. Adding a heavy-demand family narrowed the shortfall (9.9% → 9.1%), but a fixed 4% (8.2%) still does better. The reason is that ORLIB's gap is almost flat from P = 4% to 15%, while the synthetic curves are steep.

The robust lesson is the scale: for SA, a good P is a few percent of the LP bound, and the 50% that suits Gurobi costs 30–52 points of gap.

(d) With capacities, how the constraint is written matters more than P. I used the capacitated model of Li (2026, *Urban Lifeline*, Appendix A), with single sourcing:

- **The exact MILP reproduces all seven optimal values in that paper's Table 3** (433.85, 377.28, …, 349.02), which requires unrounded Euclidean distances.
- **The QUBO adds, for each facility, a slack $s_i \in [0, k]$** written with bounded binary weights (1, 2, …), and penalizes $P\cdot(\Sigma_j x_{ij} + s_i - k\cdot y_i)^2$. An exhaustive search on a small instance confirms the QUBO minimum is the CFLP optimum.
- **ORLIB's own capacities cannot be used under single sourcing:** one customer's demand, 12,912, exceeds most capacities. So I applied Appendix A's "at most k customers per facility" to the 12 ORLIB cost matrices, with k ∈ {5, 10}: 24 problems and 1,536 SA runs.

Best mean repaired gap (12 problems, seeds 0–2)	k = 5	k = 10
capacity row $\Sigma_j x_{ij} + s_i = k \cdot y_i$ (as written)	16.6%	24.4%
capacity row $\Sigma_j x_{ij} + s_i = k$ (equally exact)	**9.4%**	**7.7%**
no capacity, same problems	8.3%	8.3%



- **As written, the row locks facilities open.** When a facility's load plus slack equals k, closing it changes the row by k, which costs $k^2 \cdot P$ in one flip. SA's reads therefore open nearly every facility (15–16 of 16 on cap74, where the optimum opens 6–10), and the gap hardly depends on P.
- **Without y, the row is just as exact,** because eq. 15's $x \le y$ penalty already stops a closed facility from serving anyone. SA then behaves as in the uncapacitated case: best near P = 5–10%, and 8–18 points better than the coupled row at every $P \ge 2\%$ (paired 95% CIs over the 12 problems exclude zero).
- **The binary slack encoding has barriers of its own.**
  - With k = 10 (weights 1, 2, 4, 3), changing the slack by one unit can require two flips through states that cost 9–16·P. As a result, only 1–11% of raw reads are feasible, against 87–100% for k = 5 (weights 1, 2, 2).
  - Repair still recovers good solutions, but the raw QUBO output is rarely usable. This is a concrete instance of the paper's remark that slack variables make QUBO models harder to solve.

4. Limitations and notes

- **Classical SA only, at one setting.** Nothing here is evidence about quantum annealing, D-Wave's hybrid solver or Gurobi. The repair step is mine; the paper needed none. SA's CPU time is not comparable with the paper's time to the best solution.
- **Instances.** ORLIB's LP bound equals its optimum, so it cannot test an estimate-based rule on its own; hence the synthetic large-gap family. B-2 was not run, because UflLib was unreachable. The heavy-demand family was added *after* seeing the ORLIB failure, so ORLIB is no longer an untouched test set. The capacitated results use capacity by customer count (Appendix A, (A.4)); I checked demand-based capacity only on the Table 3 example.
- **An implementation note on Property 2.** Read literally, the index formula $\min\{|I| + m \cdot |J| + j, |I|(|J|+1)\}$ for m = i, ..., |I| clamps to the last variable when m = |I|. On the appendix example that places extra 2P entries in the x_{34} column. I built the matrix from the property's description (pairs $i < i'$ for the same customer), which reproduces the appendix matrix exactly.
- **Reproducibility.** Every run is seeded and logged. A test suite (135 tests) covers the appendix example, all 37 optima, both constructions and the *Urban Lifeline* Table 3 values. The code and data are in a GitHub repository that I can share.